

A2Q1

November 2, 2021


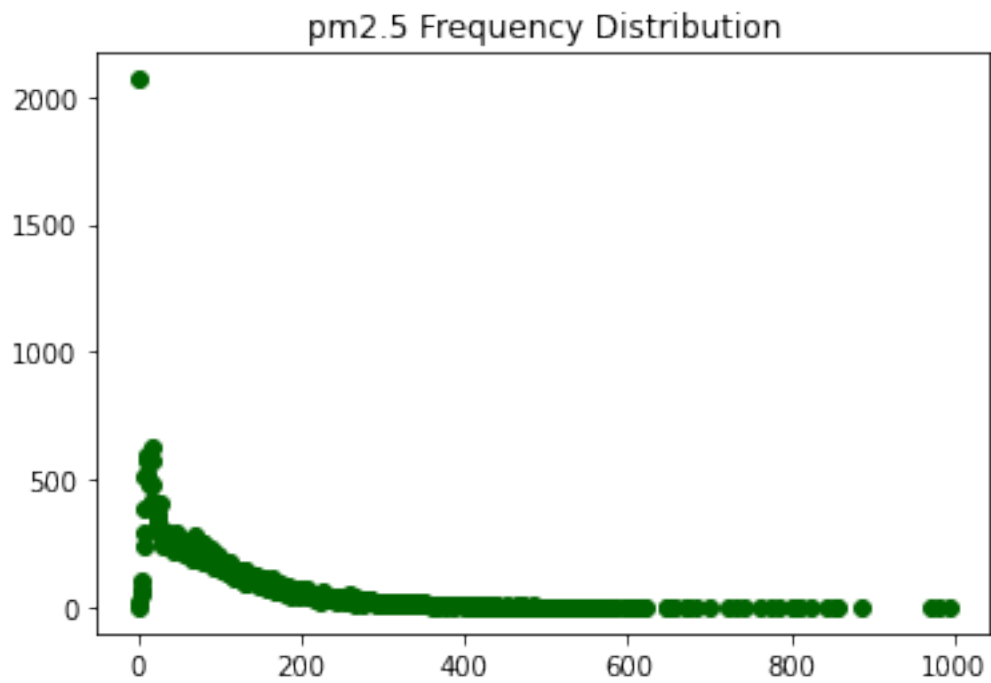
```
[4]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import sklearn
from sklearn.tree import DecisionTreeClassifier
import random
```


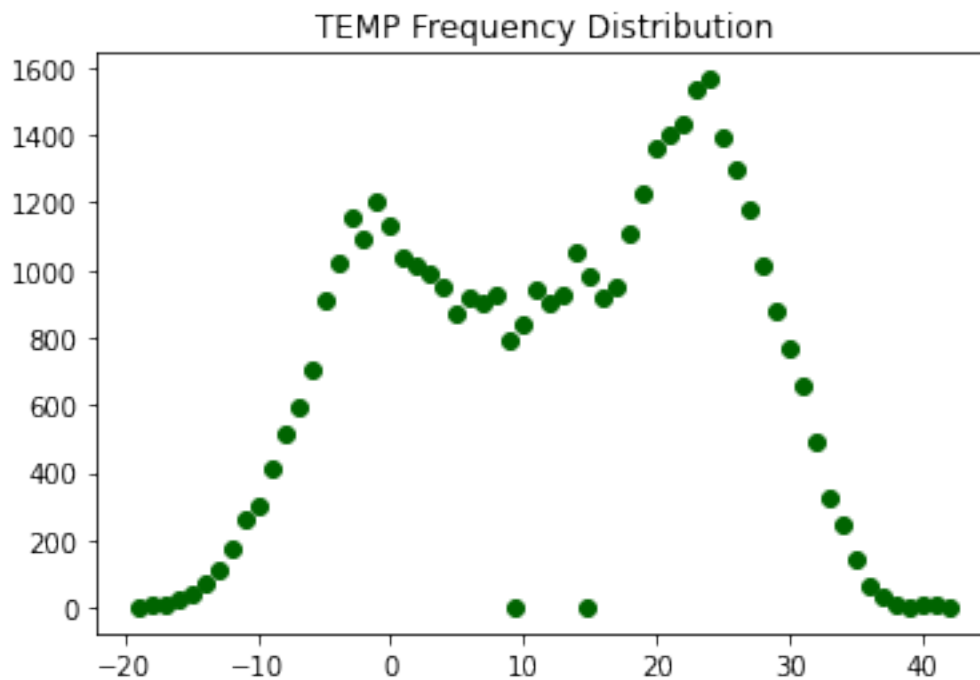
0.0.1 Data Pre - Processing and EDA

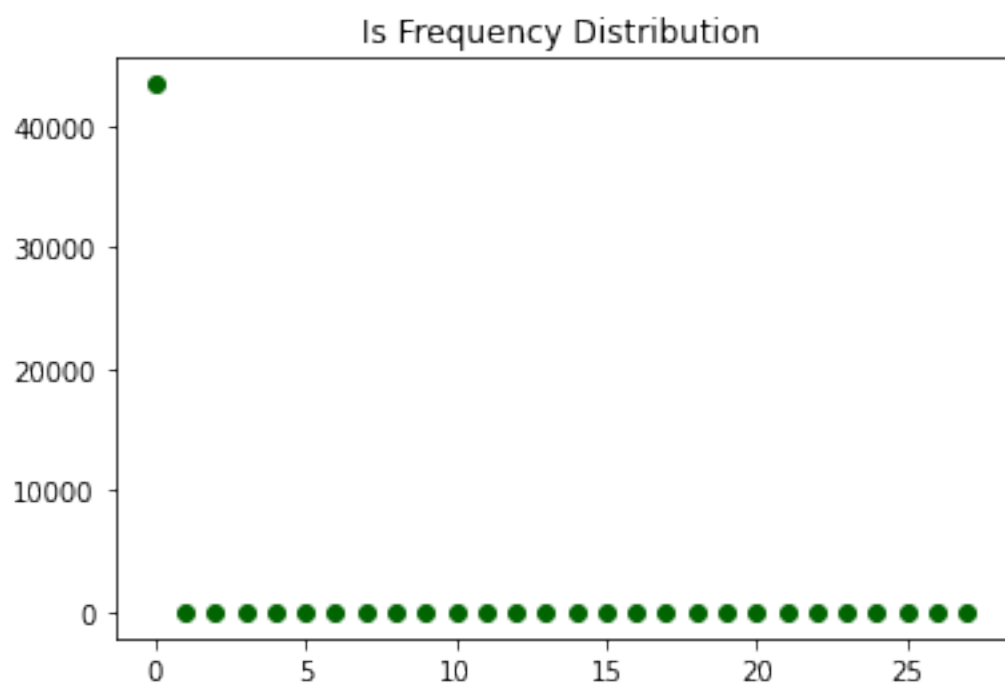
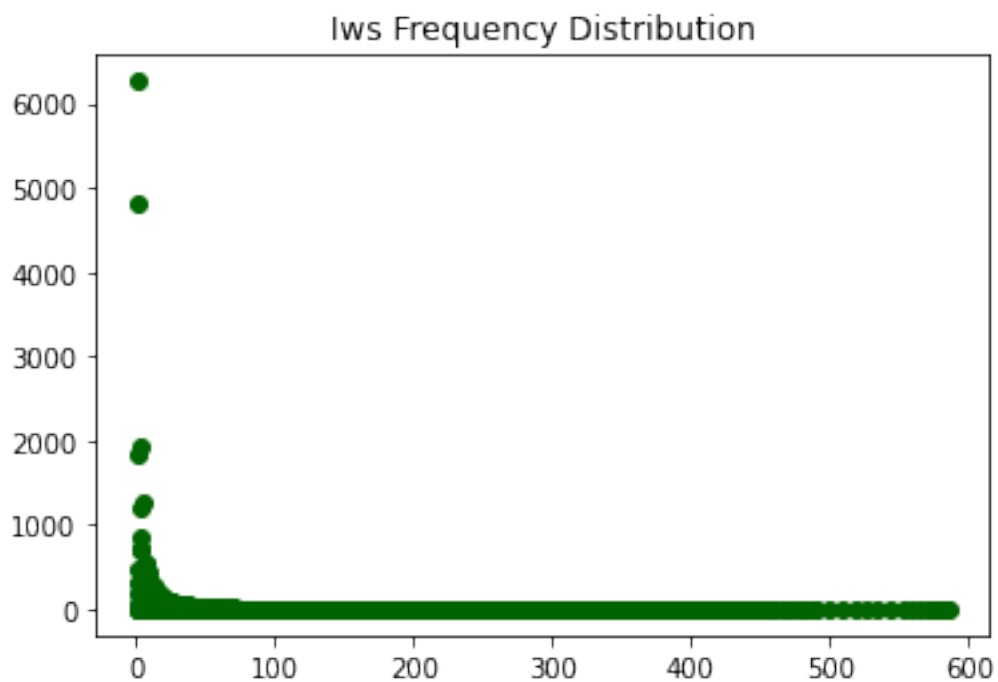
1. Replacing NaN values with 0
2. Drop the No. column as described in question
3. One Hot Encoding has been used for the 'cbwd' column
4. Plotted Frequency Distribution of Various Attributes
5. Performed Gaussian Normalization for ['DEWP', 'PRES', 'TEMP']
6. Performed Min-Max Normalization for ['pm2.5', 'Iws', 'Is', 'Ir']

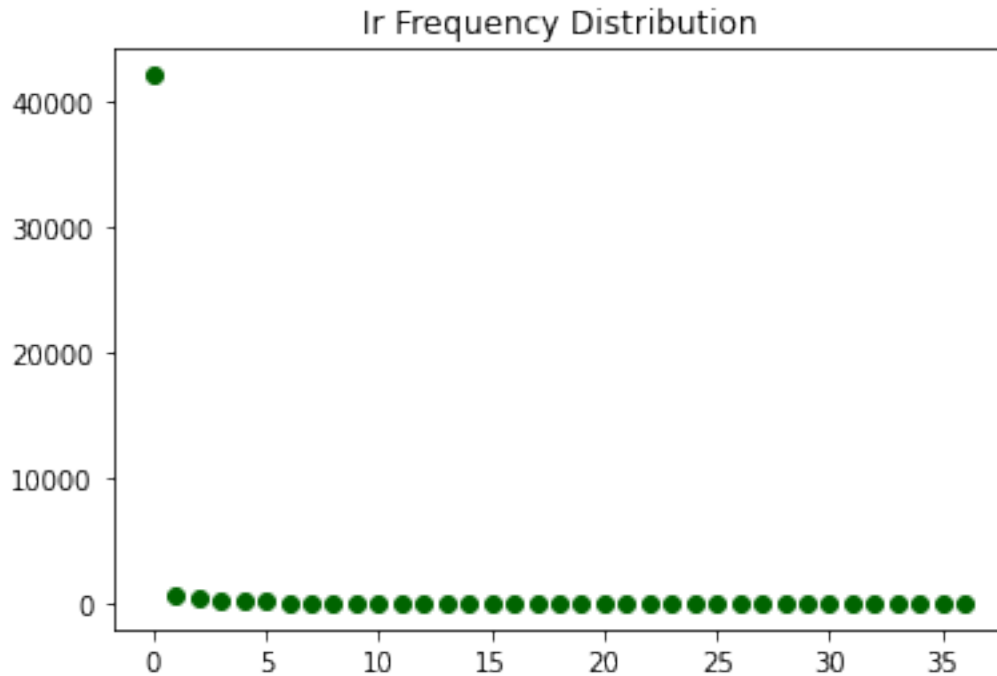
```
[6]: def eda(df):
    for column in ['pm2.5', 'DEWP', 'TEMP', 'PRES', 'Iws', 'Is', 'Ir']:
        freq = {}
        for val in df[column]:
            if not np.isnan(val):
                freq[val] = freq.get(val, 0) + 1
    x = freq.keys()
    y = freq.values()
    plt.scatter(x, y, color='darkgreen')
    plt.title(f"{column} Frequency Distribution")
    plt.show()
```

```
[7]: dir="datasetq1.csv"
df=pd.read_csv(dir)
df = df.sample(frac=1, random_state=0).reset_index(drop=True)
df.replace({np.nan:0}, inplace=True)
df=df.drop('No',axis=1)
df = pd.get_dummies(df, columns = ['cbwd'])
eda(df)
y=df["month"]
X=df.drop("month",axis=1)
```









```
[8]: gaussian_cols=['DEWP', 'PRES', 'TEMP']
minmax_cols=['pm2.5', 'Iws', 'Is', 'Ir']
for i in minmax_cols:
    X[i]=(X[i]-X[i].min())/(X[i].max()-X[i].min())
for i in gaussian_cols:
    X[i]=(X[i]-X[i].mean()/(X[i].std()))
```

```
[9]: def accuracy_score(y_pred, y):
    counter=0
    for i in range(len(y_pred)):
        if(y_pred[i]==y[i]):
            counter+=1
    return counter/len(y)

def mytraintestvalsplitt(total, valsize, testsize):
    notestrows=int(testsize*total.shape[0])
    novalrows=int(valsize*total.shape[0])
    notrainrows=total.shape[0]-notestrows-novalrows
    trainrows=total[:notrainrows]
    valrows=total[notrainrows:notrainrows+novalrows]
    testrows=total[notrainrows+novalrows:]
    return trainrows, valrows, testrows
X=X.to_numpy()
y=y.to_numpy()
```

```
X_train, X_val,X_test = mytraintestvalsplit(X, 0.15,0.15)
Y_train, Y_val,Y_test = mytraintestvalsplit(y, 0.15,0.15)
```

```
[10]: clf = DecisionTreeClassifier(criterion='gini')
clf = clf.fit(X_train, Y_train)
pred = clf.predict(X_test)
score = accuracy_score(pred, Y_test)
print(f"Accuracy with Gini comes out to be {score}")
```

Accuracy with Gini comes out to be 0.8209341244485014

```
[11]: clf = DecisionTreeClassifier(criterion='entropy')
clf = clf.fit(X_train, Y_train)
pred = clf.predict(X_test)
score = accuracy_score(pred, Y_test)
print(f"Accuracy with Entropy comes out to be {score}")
```

Accuracy with Entropy comes out to be 0.8323444393731934

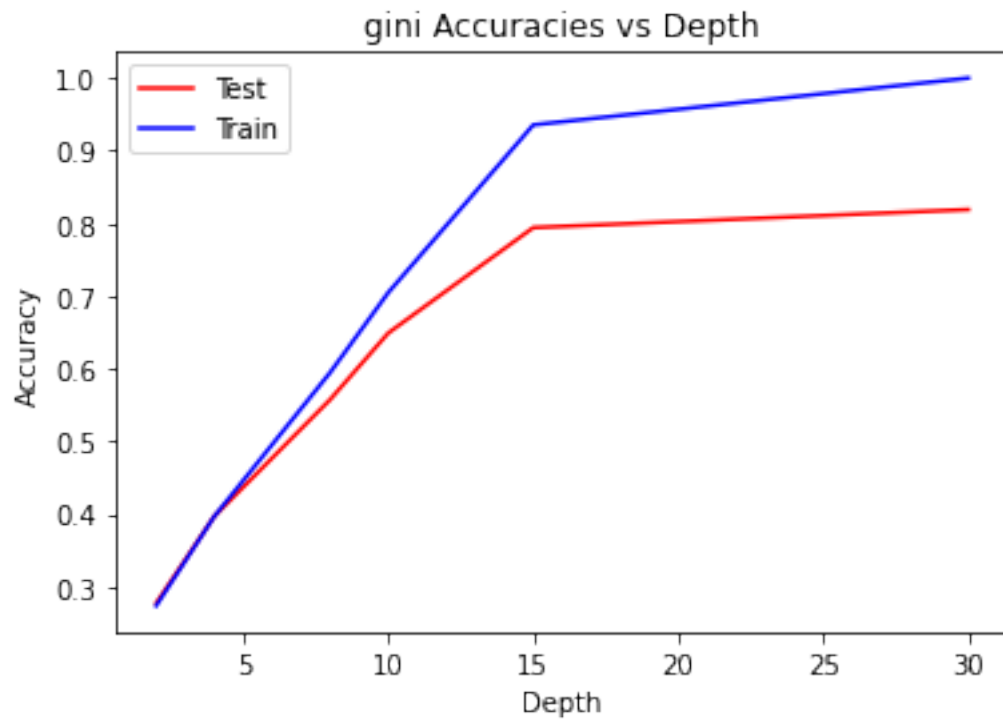
Q1 (a) Decision Tree with Entropy Gives a higher Accuracy.

```
[8]: def q1b(technique):
    max_dep=[2,4,8,10,15,30]
    score_test=[]
    score_train=[]
    for i in max_dep:
        clf = DecisionTreeClassifier(criterion=technique,max_depth=i)
        clf = clf.fit(X_train, Y_train)
        pred = clf.predict(X_test)
        score = accuracy_score(pred, Y_test)
        score_test.append(score)
        pred = clf.predict(X_train)
        score = accuracy_score(pred, Y_train)
        score_train.append(score)
    plt.plot(max_dep, score_test, label = "Test",color="red")
    plt.plot(max_dep, score_train, label = "Train",color="blue")
    plt.xlabel('Depth')
    plt.ylabel('Accuracy')
    plt.title(f'{technique} Accuracies vs Depth')

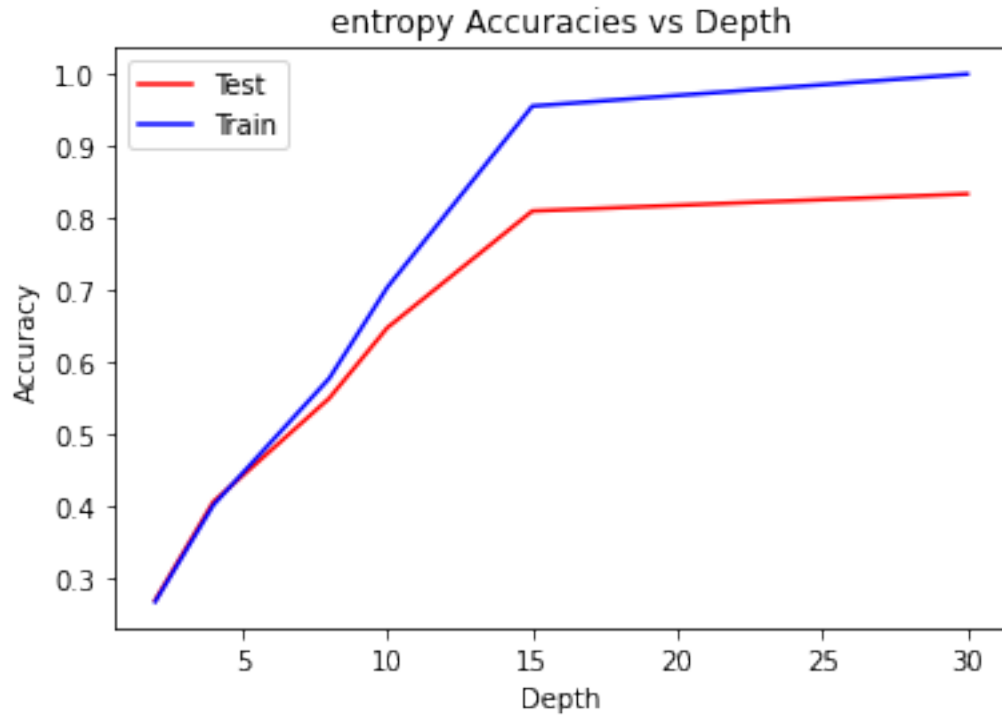
    plt.legend()
    plt.show()
    print()

    for x in range(len(score_test)):
        print(f"The accuracy score with {technique} on the test set with depth_
→{max_dep[x]} is {score_test[x]}")
```

```
q1b('gini')  
q1b('entropy')
```



The accuracy score with gini on the test set with depth 2 is 0.2771945839038491
The accuracy score with gini on the test set with depth 4 is 0.3966225467822912
The accuracy score with gini on the test set with depth 8 is 0.5575840559866119
The accuracy score with gini on the test set with depth 10 is 0.6490187129164765
The accuracy score with gini on the test set with depth 15 is 0.7941579187585578
The accuracy score with gini on the test set with depth 30 is 0.8189563365282215



The accuracy score with entropy on the test set with depth 2 is
0.27065267001369236

The accuracy score with entropy on the test set with depth 4 is
0.40696789898067853

The accuracy score with entropy on the test set with depth 8 is
0.5505857294994675

The accuracy score with entropy on the test set with depth 10 is
0.6482580252548303

The accuracy score with entropy on the test set with depth 15 is
0.8101323596531265

The accuracy score with entropy on the test set with depth 30 is
0.8338658146964856

Q1 (b) The best value of depth for both is 30. Entropy performs slightly better than Gini. (Plots shown above)

```
[9]: def q1c(nooftrees,depth,X_train,Y_train,X_test,Y_test):
    list_clf=[]

    for i in range(nooftrees):
        clf=DecisionTreeClassifier(criterion='entropy',max_depth=depth)
        temp_xtrain,_,temp_ytrain,__,=train_test_split(X_train,Y_train,test_size=.
        ↪50,shuffle=True)
```



```

clf = clf.fit(temp_xtrain, temp_ytrain)
temp_pred = clf.predict(X_test)
list_clf.append(temp_pred)

dic_pred=[]
for i in range(len(list_clf[0])):
    new_list=[]
    for j in range(len(list_clf)):
        new_list.append(list_clf[j][i])
    dic_pred.append(new_list)

def most_frequent(List):
    return max(set(List), key = List.count)

prediction=[]
for i in range(len(dic_pred)):
    prediction.append(most_frequent(dic_pred[i]))

score = accuracy_score(prediction, Y_test)
return (score)

```

```

[10]: print(f"The accuracy score for our experiment comes out to be_
↪{q1c(100,3,X_train,Y_train,X_test,Y_test)}")

```

The accuracy score for our experiment comes out to be 0.34215731020842843

This accuracy is way less than 1 (a) and (b)

This happened because the max depth of trees in the random forest were taken to be 3 which makes a decision tree with very low depth and the model is under-fitting, i.e. it is not able to use all the input features.

Q1 (d) Trying out various depths [4,8,10,15,20,30] and plotting to find the best depth.

```

[11]: max_depths=[4,8,10,15,20,30]
num_of_trees=100
training_accuracy=[]
validation_accuracy=[]
testing_accuracy=[]
for j in max_depths:
    training_accuracy.append(q1c(100,j,X_train,Y_train,X_train,Y_train))
    validation_accuracy.append(q1c(100,j,X_train,Y_train,X_val,Y_val))
    testing_accuracy.append(q1c(100,j,X_train,Y_train,X_test,Y_test))

print("Test Set Accuracies")
for x in range(len(testing_accuracy)):
    print(f"The accuracy with depth {max_depths[x]} is : {testing_accuracy[x]}")

```

```

print("Validation Set Accuracies")
for x in range(len(validation_accuracy)):
    print(f"The accuracy with depth {max_depths[x]} is :↳
    ↳{validation_accuracy[x]}")

print("Training Set Accuracies")
for x in range(len(training_accuracy)):
    print(f"The accuracy with depth {max_depths[x]} is : {training_accuracy[x]}")

plt.plot(max_depths, training_accuracy, label = "Train", color="red")
plt.plot(max_depths, validation_accuracy, label = "Validation", color="blue")
plt.plot(max_depths, testing_accuracy, label = "Test", color="green")
plt.xlabel('Max Depth')
plt.ylabel('Accuracy')
plt.title('Accuracies vs Ensemble Depth')
plt.legend()
plt.show()
print()

```

Test Set Accuracies

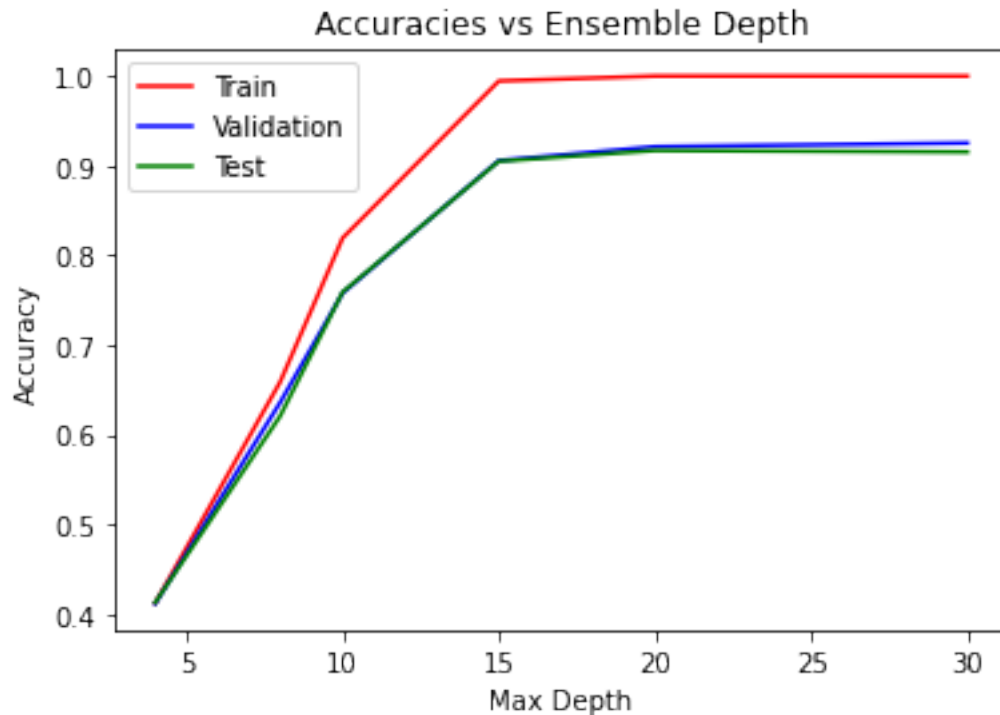
The accuracy with depth 4 is : 0.41259698767685987
 The accuracy with depth 8 is : 0.6205689943709113
 The accuracy with depth 10 is : 0.7588620112581774
 The accuracy with depth 15 is : 0.9047619047619048
 The accuracy with depth 20 is : 0.9167807698159136
 The accuracy with depth 30 is : 0.9151072569602922

Validation Set Accuracies

The accuracy with depth 4 is : 0.41137988741822606
 The accuracy with depth 8 is : 0.6357827476038339
 The accuracy with depth 10 is : 0.7577970485318728
 The accuracy with depth 15 is : 0.9055225924235509
 The accuracy with depth 20 is : 0.9210406207211319
 The accuracy with depth 30 is : 0.925148334094021

Training Set Accuracies

The accuracy with depth 4 is : 0.4119890475259143
 The accuracy with depth 8 is : 0.6586478909968055
 The accuracy with depth 10 is : 0.8192189842884151
 The accuracy with depth 15 is : 0.9942955864137166
 The accuracy with depth 20 is : 0.9997718234565487
 The accuracy with depth 30 is : 0.9997066301584198



So we can see till depth 20, the accuracy increases on all the sets(training, validation and the test set) and after that it decreases or remains nearly same. So the best depth we can use out of these will be 20. On using a higher depth we are over fitting the model and increasing the computation power required as well, so 20 turns out to be an optimum choice.

0.0.2 Now Tuning for number of Trees [100,120,140,150]

```
[18]: max_trees=[100,120,140,150]
training_accuracy=[]
validation_accuracy=[]
testing_accuracy=[]
for j in max_trees:
    training_accuracy.append(q1c(j,20,X_train,Y_train,X_train,Y_train))
    validation_accuracy.append(q1c(j,20,X_train,Y_train,X_val,Y_val))
    testing_accuracy.append(q1c(j,20,X_train,Y_train,X_test,Y_test))

print("Test Set Accuracies")
for x in range(len(testing_accuracy)):
    print(f"The accuracy with no of trees {max_trees[x]} is :_
    ↳{testing_accuracy[x]}")

print("Validation Set Accuracies")
```

```

for x in range(len(validation_accuracy)):
    print(f"The accuracy with no of trees {max_trees[x]} is :␣
    ↳{validation_accuracy[x]}")

print("Training Set Accuracies")
for x in range(len(training_accuracy)):
    print(f"The accuracy with no of trees {max_trees[x]} is :␣
    ↳{training_accuracy[x]}")

plt.plot(max_trees, training_accuracy, label = "Train", color="red")
plt.plot(max_trees, validation_accuracy, label = "Validation", color="blue")
plt.plot(max_trees, testing_accuracy, label = "Test", color="green")
plt.xlabel('Max Trees')
plt.ylabel('Accuracy')
plt.title('Accuracies vs Ensemble with different number of trees')
plt.legend()
plt.show()
print()

```

Test Set Accuracies

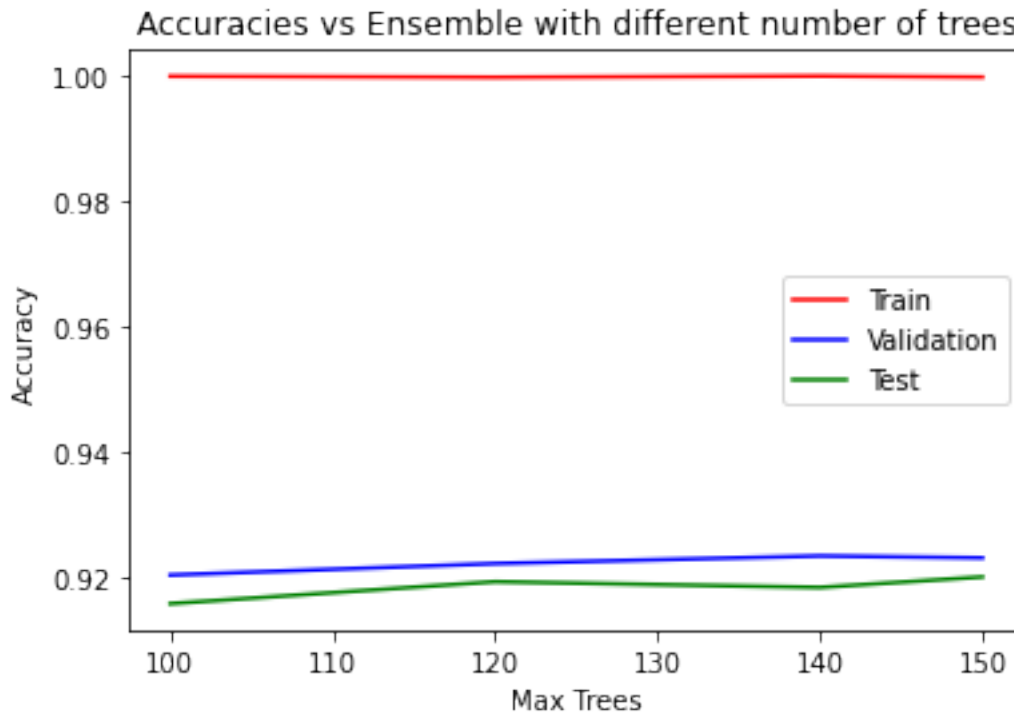
The accuracy with no of trees 100 is : 0.915715807089609
 The accuracy with no of trees 120 is : 0.9192149703331812
 The accuracy with no of trees 140 is : 0.9183021451392058
 The accuracy with no of trees 150 is : 0.9199756579948273

Validation Set Accuracies

The accuracy with no of trees 100 is : 0.9202799330594857
 The accuracy with no of trees 120 is : 0.9221055834474364
 The accuracy with no of trees 140 is : 0.9233226837060703
 The accuracy with no of trees 150 is : 0.9230184086414118

Training Set Accuracies

The accuracy with no of trees 100 is : 0.9997718234565487
 The accuracy with no of trees 120 is : 0.9996088402112263
 The accuracy with no of trees 140 is : 0.9998044201056131
 The accuracy with no of trees 150 is : 0.9996414368602907



According to the plot we can see, that the best number of trees for the ensemble out of the tested values will be 150.

Analysis : On increasing depth beyond 20, our model is overfitting and hence we see slight decrease in accuracies. On max trees=150, we are getting the best results as both the validation and test sets are giving similar results and overall accuracy is also maximum.

```
[12]: from sklearn.ensemble import AdaBoostClassifier
est=[4,8,10,15,20]
training_accuracy=[]
validation_accuracy=[]
testing_accuracy=[]
for j in est:
    clf = DecisionTreeClassifier(criterion = 'entropy',max_depth=20)
    abc = AdaBoostClassifier(base_estimator=clf, n_estimators=j, learning_rate=1,
    random_state=0)
    model = abc.fit(X_train, Y_train)
    train_pred = model.predict(X_train)
    validation_pred = model.predict(X_val)
    test_pred = model.predict(X_test)
    training_accuracy.append(accuracy_score(Y_train, train_pred))
    validation_accuracy.append(accuracy_score(Y_val, validation_pred))
    testing_accuracy.append(accuracy_score(Y_test, test_pred))
```

```

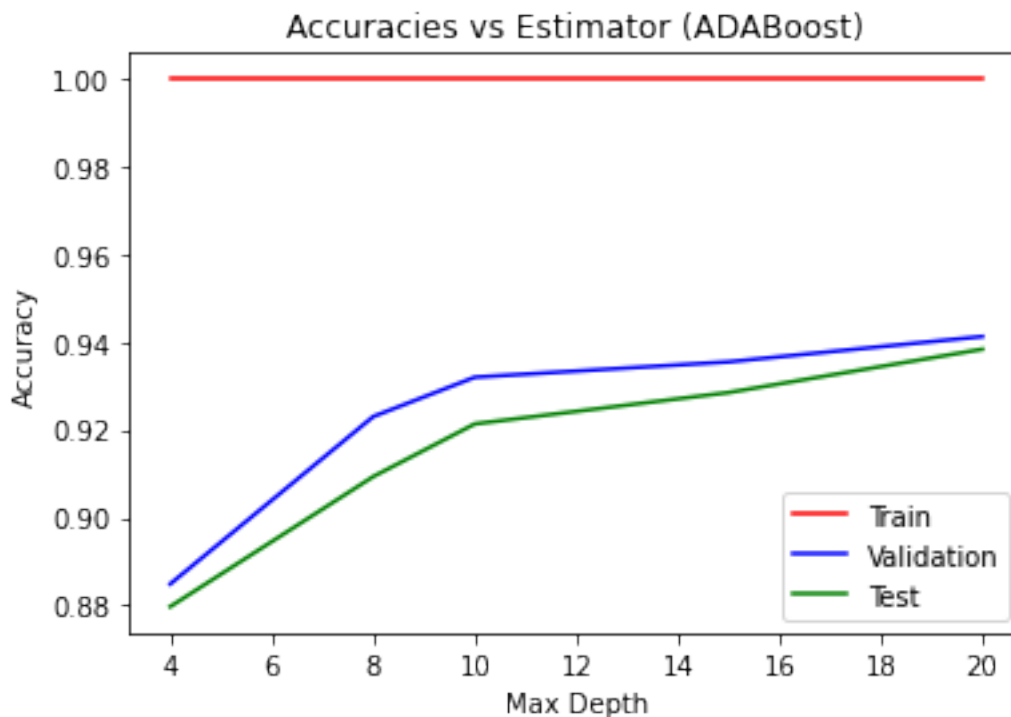
print("Test Set Accuracies")
for x in range(len(testing_accuracy)):
    print(f"The accuracy with ADABOOST with n_estimator {est[x]} is :_
    ↳{testing_accuracy[x]}")

plt.plot(est,training_accuracy,label = "Train",color="red")
plt.plot(est,validation_accuracy,label = "Validation",color="blue")
plt.plot(est,testing_accuracy,label = "Test",color="green")
plt.xlabel('Max Depth')
plt.ylabel('Accuracy')
plt.title('Accuracies vs Estimator (ADABOOST)')
plt.legend()
plt.show()

```

Test Set Accuracies

The accuracy with ADABOOST with n_estimator 4 is : 0.8796592119275826
 The accuracy with ADABOOST with n_estimator 8 is : 0.9093260307317815
 The accuracy with ADABOOST with n_estimator 10 is : 0.9213448957857904
 The accuracy with ADABOOST with n_estimator 15 is : 0.928495359805264
 The accuracy with ADABOOST with n_estimator 20 is : 0.9383842994066636



0.0.3 ADABOOST performs best for n_estimators = 20

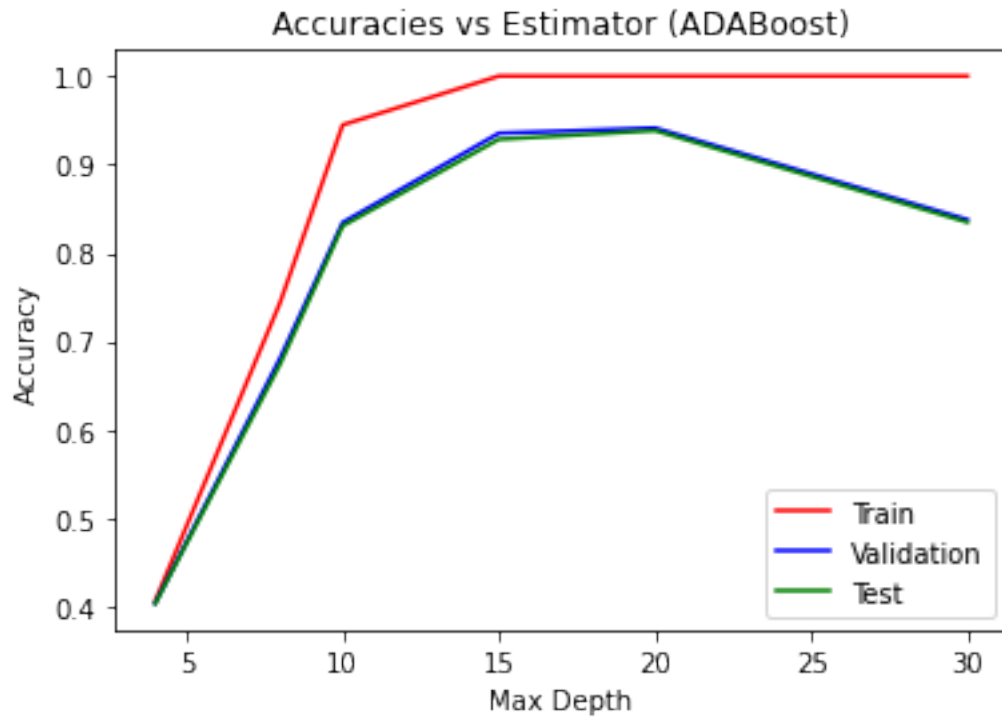
```
[13]: max_depths=[4,8,10,15,20,30]
training_accuracy=[]
validation_accuracy=[]
testing_accuracy=[]
for j in max_depths:
    clf = DecisionTreeClassifier(criterion = 'entropy', max_depth = j)
    abc = AdaBoostClassifier(base_estimator=clf, n_estimators=20,
    ↪ learning_rate=1, random_state=0)
    model = abc.fit(X_train, Y_train)
    train_pred = model.predict(X_train)
    validation_pred = model.predict(X_val)
    test_pred = model.predict(X_test)
    training_accuracy.append(accuracy_score(Y_train, train_pred))
    validation_accuracy.append(accuracy_score(Y_val, validation_pred))
    testing_accuracy.append(accuracy_score(Y_test, test_pred))

print("Test Set Accuracies")
for x in range(len(testing_accuracy)):
    print(f"The accuracy with ADABOOST with max depth of tree {max_depths[x]} is :
    ↪ {testing_accuracy[x]}")

plt.plot(max_depths, training_accuracy, label = "Train", color="red")
plt.plot(max_depths, validation_accuracy, label = "Validation", color="blue")
plt.plot(max_depths, testing_accuracy, label = "Test", color="green")
plt.xlabel('Max Depth')
plt.ylabel('Accuracy')
plt.title('Accuracies vs Estimator (ADABOOST)')
plt.legend()
plt.show()
```

Test Set Accuracies

The accuracy with ADABOOST with max depth of tree 4 is : 0.4033165982047771
The accuracy with ADABOOST with max depth of tree 8 is : 0.6744256808154572
The accuracy with ADABOOST with max depth of tree 10 is : 0.8305187889852427
The accuracy with ADABOOST with max depth of tree 15 is : 0.928495359805264
The accuracy with ADABOOST with max depth of tree 20 is : 0.9383842994066636
The accuracy with ADABOOST with max depth of tree 30 is : 0.834778639890461



Conclusion : ADABOOST with $n_estimators = 20$ is superior to RF on comparing on various decision tree depths. Our custom implementation was able to reach a max accuracy of about 92 % whereas ADABOOST with max depth = 20 and $n_estimators = 20$ has a testing accuracy of 93.8%. On comparing RF with various tree depths with ADABOOST we can see ADABOOST performs better or very similar in $max_depths=[4,8,10,15,20,30]$

[]:

Q2

November 2, 2021

```
[7]: import matplotlib.pyplot as plt
class NN():
    """
    Neural Network Classifier
    """
    def __init__(self, num_layers, layer_size, activation, lr, weightinit,
    ↪batch_size, epochs):
        self.num_layers, self.layer_size, self.activation, self.lr, self.
    ↪weightinit, self.batch_size, self.epochs = num_layers, layer_size,
    ↪activation, lr, weightinit, batch_size, epochs

    def activationfn(self,X):
        string=f"self.{self.activation}(X)"
        return eval(string)

    def gradfn(self,X):
        string=f"self.{self.activation}_grad(X)"
        return eval(string)

    def score(self, X, y):
        y_pred = self.predict(X)
        counter=0
        for i in range(len(y_pred)):
            if(y_pred[i]==y):
                counter+=1
        return counter/len(y)

    def initialization(self):
        """
        We make matrices of (layer_size[i], layer_size[i+1]) as we need
        Wji as the params, all combinations that are possible.
        """

        params = {}
        mylayers = self.layer_size
```

```

for i in range(0,self.num_layers-1):
    params["b" + str(i+1)] = np.zeros((1,mylayers[i+1]))

    if(self.weightinit == 'normal'):
        thislayer = np.random.normal(size = (mylayers[i],mylayers[i+1]))*0.01

    elif(self.weightinit == 'zero'):
        thislayer = np.zeros((mylayers[i],mylayers[i+1]))

    else:
        thislayer = np.random.rand(mylayers[i],mylayers[i+1])*0.01

    params["W" + str(i+1)] = thislayer

self.params = params

def crossentropyloss(self, Amatrix, Y):
    temp=Amatrix[np.arange(len(Y)), Y.argmax(axis=1)]
    temp=np.where(temp>0.0000000000000001,temp,0.0000000000000001)
    logp = - np.log(temp)
    celoss = np.sum(logp)/len(Y)
    return celoss

def forward_prop(self,X,params):

    A = X
    myactivations = {}
    before_activation = {}
    numhiddenlayers=self.num_layers-2
    for i in range(numhiddenlayers):
        A_prev = A
        Z = np.dot(A_prev, params["W" + str(i+1)]) + params["b" + str(i+1)]
        before_activation["Z" + str(i+1)] = Z
        A=self.activationfn(Z)
        myactivations["A" + str(i+1)] = A
        A_prev = A

    ZL = np.dot(A_prev, params["W" + str(numhiddenlayers+1)]) + params["b" +
↪str(numhiddenlayers+1)]
    AL = (np.exp(ZL)/(np.sum(np.exp(ZL),axis = 1, keepdims = True))) #↪
↪SoftMax
    myactivations["A" + str(numhiddenlayers+1)] = AL
    before_activation["Z" + str(numhiddenlayers+1)] = ZL

```

```

    return AL, myactivations, before_activation
def backward_prop(self, X, Y, before_activation, myactivations):
    """
    Complete Backward Prop
    """
    gradients = {}
    Lay = self.num_layers-1
    myactivations["A0"] = X

    A = myactivations["A" + str(Lay)]
    dZ = A - Y

    dW = np.dot(myactivations["A" + str(Lay-1)].T, dZ)/len(X)
    db = np.sum(dZ, axis=0, keepdims=True) / len(X)
    dAp = np.dot(dZ, self.params["W" + str(Lay)].T)

    gradients["db" + str(Lay)] = db
    gradients["dW" + str(Lay)] = dW

    for l in range(Lay - 1, 0, -1):
        dGrad=self.gradfn(before_activation["Z" + str(l)])
        dZ = dAp * dGrad
        dW = (1/len(X)) * np.dot(myactivations["A" + str(l - 1)].T, dZ)
        db = (1/len(X)) * np.sum(dZ, axis=0, keepdims=True)
        if l > 1:
            dAp = np.dot(dZ, self.params["W" + str(l)].T)
            gradients["dW" + str(l)] = dW
            gradients["db" + str(l)] = db

    #Update the params
    for i in range(Lay):
        self.params["W" + str(i+1)] -= self.lr*gradients["dW" + str(i+1)]
        self.params["b" + str(i+1)] -= self.lr*gradients["db" + str(i+1)]

def fit(self, X, y, x_val, y_val):
    """
    Train the model.
    """
    self.initialization()
    self.classes=int(np.max(y))
    m=X.shape[0]
    y = self.converttoprobvsclassmatrix(y)

```

```

train_losses = []
val_losses = []

noofbatches=m//self.batch_size
combined_train_data=[]

for i in range(0,noofbatches):
    myX=X[self.batch_size*i:self.batch_size*(i+1),:]
    myY=y[self.batch_size*i:self.batch_size*(i+1),:]
    combined_train_data.append((myX,myY))

for epoch in range(self.epochs):
    print(f"Epoch: {epoch+1} ", end=' ' )
    trainbatchloss = []
    valbatchloss = []

    for batch_x, batch_y in combined_train_data:
        A, activations, preactivations = self.forward_prop(batch_x,self.
→params)
        train_cost = self.crossentropyloss(A,batch_y)
        trainbatchloss.append(train_cost)
        self.backward_prop(batch_x,batch_y,preactivations, activations )
        proba = self.predict_proba(x_val)
        valloss = self.crossentropyloss(proba, self.
→converttoprobvsclassmatrix(y_val))
        valbatchloss.append(valloss)

    train_losses.append(np.array(trainbatchloss).mean())
    val_losses.append(np.array(valbatchloss).mean())

self.train_losses = train_losses
self.val_losses = val_losses

def converttoprobvsclassmatrix(self, y):
    m = len(y)
    myy = np.zeros((int(m),self.classes+1))
    for i in range(m):
        l = int(y[i])

```

```

        myy[i,1] = 1
    return myy
def predict_proba(self, inp_X):
    proba,temp,temp2 = self.forward_prop(inp_X,self.params)
    return proba
def predict(self, X):
    proba = self.predict_proba(X)
    y_pred = np.argmax(proba, axis = 1)
    return y_pred

def score(self, X, y):
    y_pred = self.predict(X)
    counter=0

    for i in range(len(y_pred)):
        if(y_pred[i]==y[i]):
            counter+=1
    return counter/len(y)

def relu(self, X):
    return X * (X>=0)

def relu_grad(self, X):
    return 1*(X>=0)

def leakyrelu(self, X):
    return np.where(X > 0, X, X * 0.01)
def leakyrelu_grad(self,z):
    f = np.maximum(z, 0.01*z)
    grad = np.where(f>0, 1, 0.01)
    return grad

def sigmoid(self, X):
    return 1/(1+np.exp(-X))

def sigmoid_grad(self, X):
    return self.sigmoid(X) *(1-self.sigmoid (X))

def linear(self, X):
    return X

def linear_grad(self, X):
    return np.ones(X.shape)

def tanh(self, X):
    return np.tanh(X)

```

```

def tanh_grad(self, X):
    return 1-self.tanh(X)*self.tanh(X)

def softmax(self, X):
    exp = np.exp(X)
    return exp/(np.sum(exp,axis = 1, keepdims = True))

def softmax_grad(self, X):
    return self.softmax(X) *(1-self.softmax(X))

```

Data Preprocessing

1. Converted the 28x28 images to 784 features for every image.
2. Tried Normalization - Both Min-Max and Gaussian Normalization i tried but they gave worse results.

```

[2]: import numpy as np
import idx2numpy

dir="train-images.idx3-ubyte"
train_X = idx2numpy.convert_from_file(dir)
dir="train-labels.idx1-ubyte"
train_Y = idx2numpy.convert_from_file(dir)
dir="t10k-images.idx3-ubyte"
test_X = idx2numpy.convert_from_file(dir)
dir="t10k-labels.idx1-ubyte"
test_Y = idx2numpy.convert_from_file(dir)

```

```

[3]: import pandas as pd
total_data_X=[]
total_data_Y=[]
for i in range(len(train_Y)):
    mylist=list(train_X[i].ravel())
    total_data_X.append(mylist)
    total_data_Y.append(train_Y[i])
for i in range(len(test_Y)):
    mylist=list(test_X[i].ravel())
    total_data_X.append(mylist)
    total_data_Y.append(test_Y[i])

total_data_X=np.array(total_data_X)
#total_data_X = total_data_X/255
total_data_Y=np.array(total_data_Y)

```

```
def mytraintestvalsplitt(total, valsize, testsize):
    notestrows=int(testsize*total.shape[0])
    novalrows=int(valsize*total.shape[0])
    notrainrows=total.shape[0]-notestrows-novalrows
    trainrows=total[:notrainrows]
    valrows=total[notrainrows:notrainrows+novalrows]
    testrows=total[notrainrows+novalrows:]
    return trainrows, valrows, testrows
```

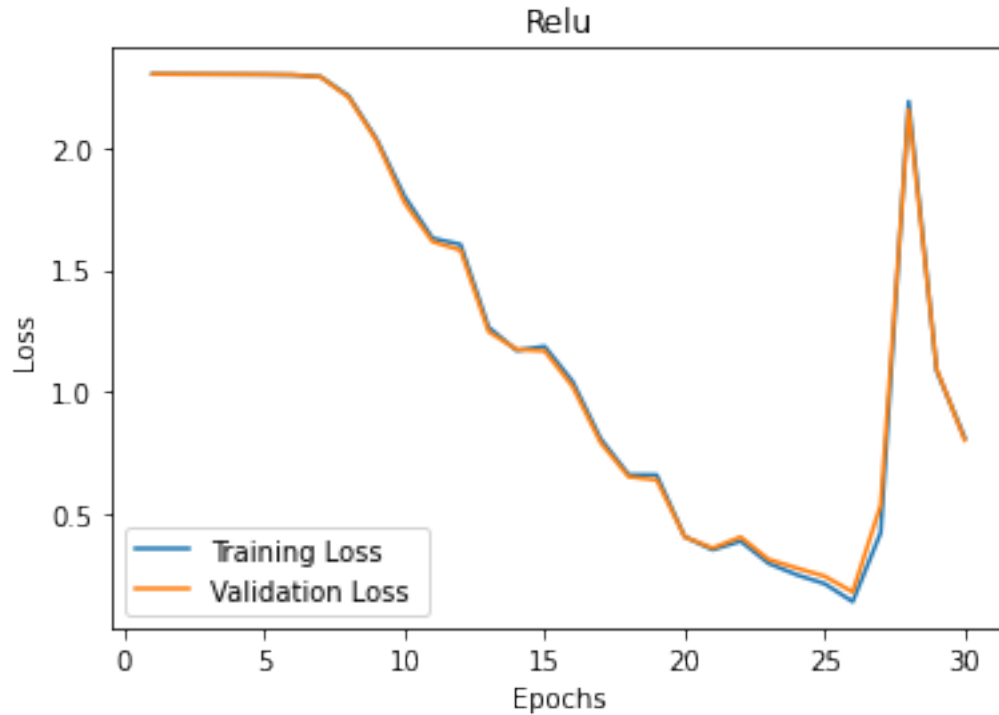
```
X_train, X_val, X_test = mytraintestvalsplitt(total_data_X, 0.2, 0.1)
Y_train, Y_val, Y_test = mytraintestvalsplitt(total_data_Y, 0.2, 0.1)
```

#Min Max Scaling#

```
[4]: nn = NN(6, [784, 256, 128, 64, 32, 10], 'relu', 0.08, 'normal', len(X_train)//20,
    ↪ 30)
nn.fit(X_train, Y_train, X_val, Y_val)
print()
print("The Accuracy Score for Relu Activation is with normal Initialization is")
print(nn.score(X_test, Y_test))
```

```
Epoch: 1 Epoch: 2 Epoch: 3 Epoch: 4 Epoch: 5 Epoch: 6 Epoch: 7 Epoch: 8 Epoch: 9
Epoch: 10 Epoch: 11 Epoch: 12 Epoch: 13 Epoch: 14 Epoch: 15 Epoch: 16 Epoch: 17
Epoch: 18 Epoch: 19 Epoch: 20 Epoch: 21 Epoch: 22 Epoch: 23 Epoch: 24 Epoch: 25
Epoch: 26 Epoch: 27 Epoch: 28 Epoch: 29 Epoch: 30
The Accuracy Score for Relu Activation is with normal Initialization is
0.8475714285714285
```

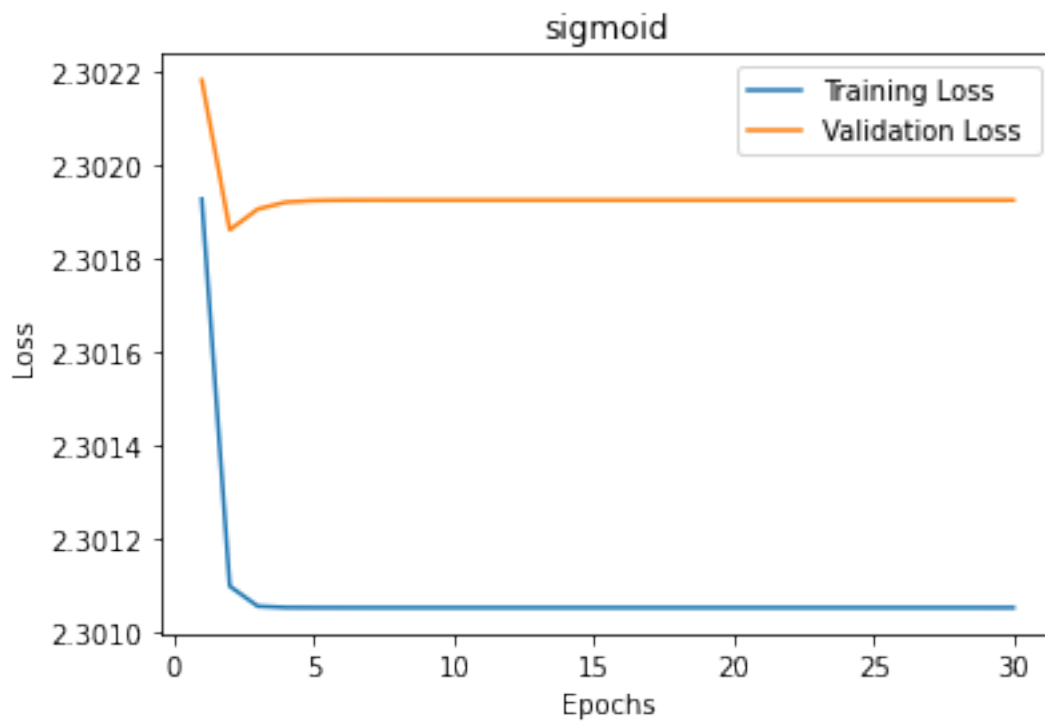
```
[5]: alltheweights={}
alltheweights['relu'] = nn.params
plt.plot(list(range(1, len(nn.train_losses) + 1)), nn.train_losses, label =
    ↪ "Training Loss ")
plt.plot(list(range(1, len(nn.val_losses) + 1)), nn.val_losses, label =
    ↪ "Validation Loss ")
plt.ylabel('Loss')
plt.legend()
plt.xlabel('Epochs')
plt.title("Relu")
plt.show()
```




```
[6]: alltheactivations=['sigmoid','tanh','linear','leakyrelu', 'softmax']

for x in alltheactivations:
    nn = NN(6, [784,256, 128, 64, 32,10], x, 0.08, 'normal', len(X_train)//20, 30)
    nn.fit(X_train,Y_train,X_val,Y_val)
    print(f"The Accuracy Score for {x} activation fn is :")
    print(nn.score(X_test,Y_test))
    alltheweights[x] = nn.params
    plt.plot(list(range(1,len(nn.train_losses) + 1)),nn.train_losses, label =_
    ↪"Training Loss " )
    plt.plot(list(range(1,len(nn.val_losses) + 1)),nn.val_losses, label =_
    ↪"Validation Loss " )
    plt.ylabel('Loss')
    plt.legend()
    plt.xlabel('Epochs')
    plt.title(x)
    plt.show()
```

```
Epoch: 1 Epoch: 2 Epoch: 3 Epoch: 4 Epoch: 5 Epoch: 6 Epoch: 7 Epoch: 8 Epoch: 9
Epoch: 10 Epoch: 11 Epoch: 12 Epoch: 13 Epoch: 14 Epoch: 15 Epoch: 16 Epoch: 17
Epoch: 18 Epoch: 19 Epoch: 20 Epoch: 21 Epoch: 22 Epoch: 23 Epoch: 24 Epoch: 25
Epoch: 26 Epoch: 27 Epoch: 28 Epoch: 29 Epoch: 30 The Accuracy Score for sigmoid
activation fn is :
0.11357142857142857
```

Epoch: 1 Epoch: 2 Epoch: 3 Epoch: 4 Epoch: 5 Epoch: 6 Epoch: 7 Epoch: 8 Epoch: 9
Epoch: 10 Epoch: 11 Epoch: 12 Epoch: 13 Epoch: 14 Epoch: 15 Epoch: 16 Epoch: 17
Epoch: 18 Epoch: 19 Epoch: 20 Epoch: 21 Epoch: 22 Epoch: 23 Epoch: 24 Epoch: 25
Epoch: 26 Epoch: 27 Epoch: 28 Epoch: 29 Epoch: 30 The Accuracy Score for tanh
activation fn is :
0.11357142857142857



Epoch: 1 Epoch: 2 Epoch: 3

C:\Users\Pritish\AppData\Local\Temp\ipykernel_14372\609547596.py:76:

RuntimeWarning: overflow encountered in exp

AL = (np.exp(ZL)/(np.sum(np.exp(ZL),axis = 1, keepdims = True))) # SoftMax

C:\Users\Pritish\AppData\Local\Temp\ipykernel_14372\609547596.py:76:

RuntimeWarning: invalid value encountered in true_divide

AL = (np.exp(ZL)/(np.sum(np.exp(ZL),axis = 1, keepdims = True))) # SoftMax

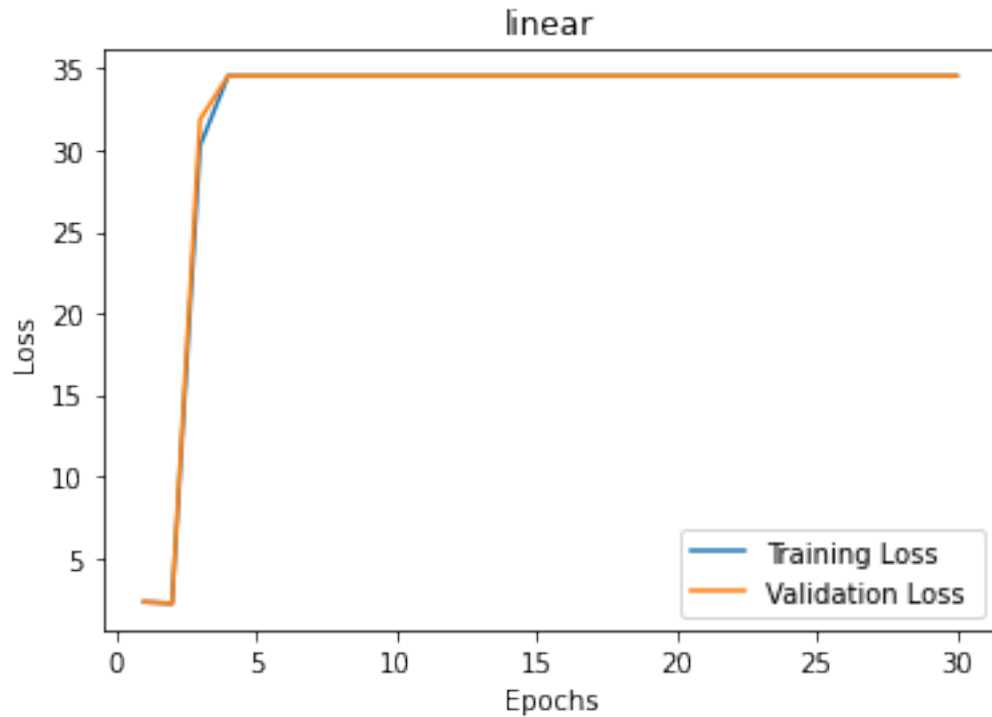
Epoch: 4 Epoch: 5 Epoch: 6 Epoch: 7 Epoch: 8 Epoch: 9 Epoch: 10 Epoch: 11 Epoch:

12 Epoch: 13 Epoch: 14 Epoch: 15 Epoch: 16 Epoch: 17 Epoch: 18 Epoch: 19 Epoch:

20 Epoch: 21 Epoch: 22 Epoch: 23 Epoch: 24 Epoch: 25 Epoch: 26 Epoch: 27 Epoch:

28 Epoch: 29 Epoch: 30 The Accuracy Score for linear activation fn is :

0.10128571428571428



Epoch: 1

```
-----
IndexError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_14372\252400328.py in <module>
      3 for x in alltheactivations:
      4     nn = NN(6, [784,256, 128, 64, 32,10], x, 0.08, 'normal', len(X_train), /
↪20, 30)
----> 5     nn.fit(X_train,Y_train,X_val,Y_val)
      6     print(f"The Accuracy Score for {x} activation fn is :")
      7     print(nn.score(X_test,Y_test))

~\AppData\Local\Temp\ipykernel_14372\609547596.py in fit(self, X, y, x_val,
↪y_val)
    149         train_cost = self.crossentropyloss(A,batch_y)
    150         trainbatchloss.append(train_cost)
--> 151         self.backward_prop(batch_x,batch_y,preactivations,
↪activations )
    152         proba = self.predict_proba(x_val)
    153         valloss = self.crossentropyloss(proba, self.
↪converttoprobvsclassmatrix(y_val))

~\AppData\Local\Temp\ipykernel_14372\609547596.py in backward_prop(self, X, Y,
↪before_activation, myactivations)
```

```

    99
   100     for l in range(Lay - 1, 0, -1):
--> 101         dGrad=self.gradfn(before_activation["Z" + str(l)])
   102         dZ = dAp * dGrad
   103         dW = (1/len(X)) * np.dot(myactivations["A" + str(l - 1)].T, d)

~\AppData\Local\Temp\ipykernel_14372\609547596.py in gradfn(self, X)
    13     def gradfn(self,X):
    14         string=f"self.{self.activation}_grad(X)"
---> 15     return eval(string)
    16
    17     def score(self, X, y):

<string> in <module>

~\AppData\Local\Temp\ipykernel_14372\609547596.py in leakyrelu_grad(self, X)
    200     def leakyrelu_grad(self,X):
    201         dx = np.ones_like(x)
--> 202         dx[X < 0] = 0.01
    203         return dx
    204

IndexError: too many indices for array: array is 0-dimensional, but 2 were
↳indexed

```

[8]: *#Re ran leaky relu because of error*

```


alltheactivations=['leakyrelu', 'softmax']

for x in alltheactivations:
    nn = NN(6, [784,256, 128, 64, 32,10], x, 0.08, 'normal', len(X_train)//20, 30)
    nn.fit(X_train,Y_train,X_val,Y_val)
    print(f"The Accuracy Score for {x} activation fn is :")
    print(nn.score(X_test,Y_test))
    alltheweights[x] = nn.params
    plt.plot(list(range(1,len(nn.train_losses) + 1)),nn.train_losses, label =↳
↳"Training Loss " )
    plt.plot(list(range(1,len(nn.val_losses) + 1)),nn.val_losses, label =↳
↳"Validation Loss " )
    plt.ylabel('Loss')
    plt.legend()
    plt.xlabel('Epochs')
    plt.title(x)
    plt.show()

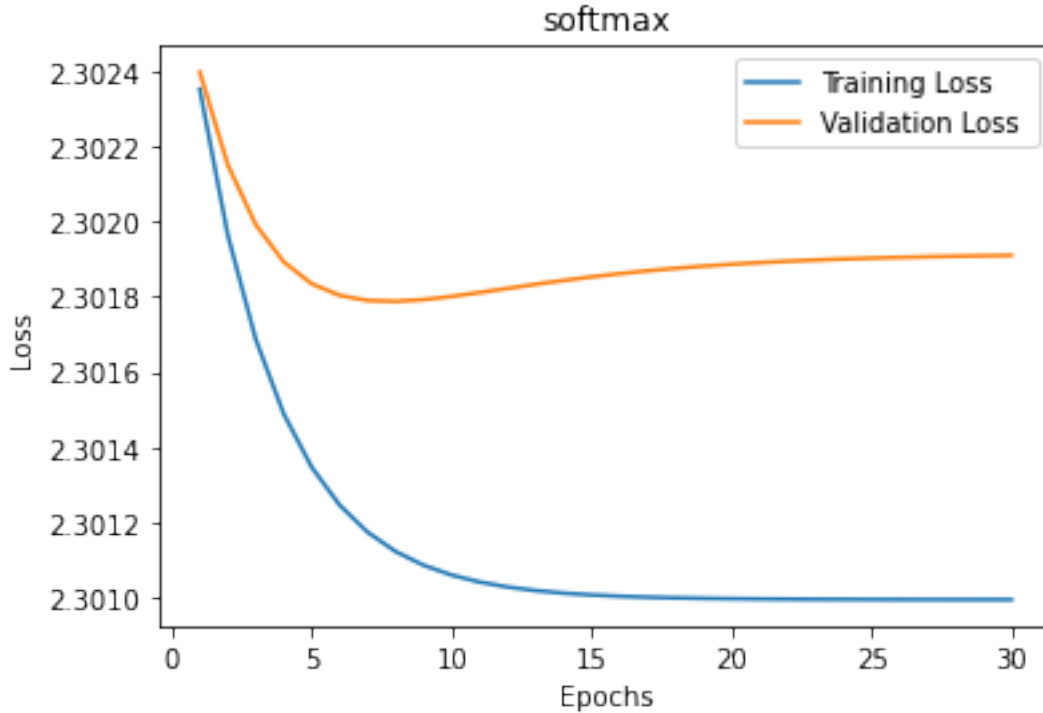
```

Epoch: 1 Epoch: 2 Epoch: 3 Epoch: 4 Epoch: 5 Epoch: 6 Epoch: 7 Epoch: 8 Epoch: 9
Epoch: 10 Epoch: 11 Epoch: 12 Epoch: 13 Epoch: 14 Epoch: 15 Epoch: 16 Epoch: 17

Epoch: 18 Epoch: 19 Epoch: 20 Epoch: 21 Epoch: 22 Epoch: 23 Epoch: 24 Epoch: 25
Epoch: 26 Epoch: 27 Epoch: 28 Epoch: 29 Epoch: 30 The Accuracy Score for
leakyrelu activation fn is :
0.7015714285714286



Epoch: 1 Epoch: 2 Epoch: 3 Epoch: 4 Epoch: 5 Epoch: 6 Epoch: 7 Epoch: 8 Epoch: 9
Epoch: 10 Epoch: 11 Epoch: 12 Epoch: 13 Epoch: 14 Epoch: 15 Epoch: 16 Epoch: 17
Epoch: 18 Epoch: 19 Epoch: 20 Epoch: 21 Epoch: 22 Epoch: 23 Epoch: 24 Epoch: 25
Epoch: 26 Epoch: 27 Epoch: 28 Epoch: 29 Epoch: 30 The Accuracy Score for softmax
activation fn is :
0.11357142857142857



- (b) From the above test accuracies, it is quite evident that “ReLU” works best. It has the best test accuracy.

The worst performance is shown in the case of the “linear” activation function as the accuracy is least in that case. It does not perform any activation and is suitable for single layers only, in our multi-layer case it does not perform well.

For “tanh”, also we weren’t able to make a good model. Maybe with more epochs we could get a desired result, but since i was low on computation power. I used only 30 epochs for these experiments.

For “relu”, it shows the best behaviour it reacts very fast to the changes and to the data. The only noticeable problem is that after certain number of epochs we tend to go away from the global

For “linear”, the training loss keeps in decreasing while validation loss fluctuates a bit and then stablizes.

For “sigmoid”, it was stuck at a local minima for a long time and the moment it got out, the training loss started decreasing while validation loss shot and then decreased to a nominal value.

- (c) In all cases the output layer should have an activation function of “softmax” with the number of nodes = the number of classes i.e. 10 (0-9), as it converts the outputs to probabilities out of which the highest is taken to get the predicted label.

```
[ ]: for activation in ["logistic", "tanh", "identity", "relu"]:
    nn = MLPClassifier(activation=activation, hidden_layer_sizes=[256, 128, 64,
→32], learning_rate_init=0.08, max_iter=30, solver="sgd", alpha = 0)
    nn.fit(X_train, Y_train)
    print(f'Test accuracy for {activation} = {nn.score(X_test, Y_test)}')
```


Test accuracy for logistic = 0.814 Test accuracy for tanh = 0.103 Test accuracy for identity = 0
Test accuracy for relu = 0.869523809523809

(d) In case of sigmoid (logistic) sklearn implementation is far superior. In the case of “linear”, the accuracy on test in custom implementation comes out to be better than the one in sklearn’s as it did not converge with the given parameters. In case of tanh we get similar results as 30 epochs are not sufficient which is what we have taken for our experiments. Sklearn also performs best for relu just like ours and the accuracies are also closeby.


```
[9]: #Bonus
lrs=[0.001,0.01,0.1,1]

for x in lrs:
    nn = NN(6, [784,256, 128, 64, 32,10], 'relu', x, 'normal', len(X_train)//20,
→30)
    nn.fit(X_train,Y_train,X_val,Y_val)
    print()
    print(f"The Accuracy Score for {x} learning rate is :")
    print(nn.score(X_test,Y_test))
    alltheweights[x] = nn.params
    plt.plot(list(range(1,len(nn.train_losses) + 1)),nn.train_losses, label =
→"Training Loss " )
    plt.plot(list(range(1,len(nn.val_losses) + 1)),nn.val_losses, label =
→"Validation Loss " )
    plt.ylabel('Loss')
    plt.legend()
    plt.xlabel('Epochs')
    plt.title(x)
    plt.show()
```


Epoch: 1 Epoch: 2 Epoch: 3 Epoch: 4 Epoch: 5 Epoch: 6 Epoch: 7 Epoch: 8 Epoch: 9
Epoch: 10 Epoch: 11 Epoch: 12 Epoch: 13 Epoch: 14 Epoch: 15 Epoch: 16 Epoch: 17
Epoch: 18 Epoch: 19 Epoch: 20 Epoch: 21 Epoch: 22 Epoch: 23 Epoch: 24 Epoch: 25
Epoch: 26 Epoch: 27 Epoch: 28 Epoch: 29 Epoch: 30
The Accuracy Score for 0.001 learning rate is :
0.11357142857142857



Epoch: 1 Epoch: 2 Epoch: 3 Epoch: 4 Epoch: 5 Epoch: 6 Epoch: 7 Epoch: 8 Epoch: 9
 Epoch: 10 Epoch: 11 Epoch: 12 Epoch: 13 Epoch: 14 Epoch: 15 Epoch: 16 Epoch: 17
 Epoch: 18 Epoch: 19 Epoch: 20 Epoch: 21 Epoch: 22 Epoch: 23 Epoch: 24 Epoch: 25
 Epoch: 26 Epoch: 27 Epoch: 28 Epoch: 29 Epoch: 30
 The Accuracy Score for 0.01 learning rate is :
 0.11357142857142857



Epoch: 1 Epoch: 2 Epoch: 3 Epoch: 4 Epoch: 5 Epoch: 6 Epoch: 7 Epoch: 8 Epoch: 9
 Epoch: 10 Epoch: 11 Epoch: 12 Epoch: 13 Epoch: 14 Epoch: 15 Epoch: 16 Epoch: 17
 Epoch: 18 Epoch: 19 Epoch: 20 Epoch: 21 Epoch: 22 Epoch: 23 Epoch: 24 Epoch: 25
 Epoch: 26 Epoch: 27 Epoch: 28 Epoch: 29 Epoch: 30
 The Accuracy Score for 0.1 learning rate is :
 0.6661428571428571



Epoch: 1

C:\Users\Pritish\AppData\Local\Temp\ipykernel_14372\692279156.py:76:

RuntimeWarning: overflow encountered in exp

```
AL = (np.exp(ZL)/(np.sum(np.exp(ZL),axis = 1, keepdims = True))) # SoftMax
```

C:\Users\Pritish\AppData\Local\Temp\ipykernel_14372\692279156.py:76:


RuntimeWarning: invalid value encountered in true_divide

```
AL = (np.exp(ZL)/(np.sum(np.exp(ZL),axis = 1, keepdims = True))) # SoftMax
```

Epoch: 2 Epoch: 3 Epoch: 4 Epoch: 5 Epoch: 6 Epoch: 7 Epoch: 8 Epoch: 9 Epoch:
10 Epoch: 11 Epoch: 12 Epoch: 13 Epoch: 14 Epoch: 15 Epoch: 16 Epoch: 17 Epoch:
18 Epoch: 19 Epoch: 20 Epoch: 21 Epoch: 22 Epoch: 23 Epoch: 24 Epoch: 25 Epoch:
26 Epoch: 27 Epoch: 28 Epoch: 29 Epoch: 30

The Accuracy Score for 1 learning rate is :

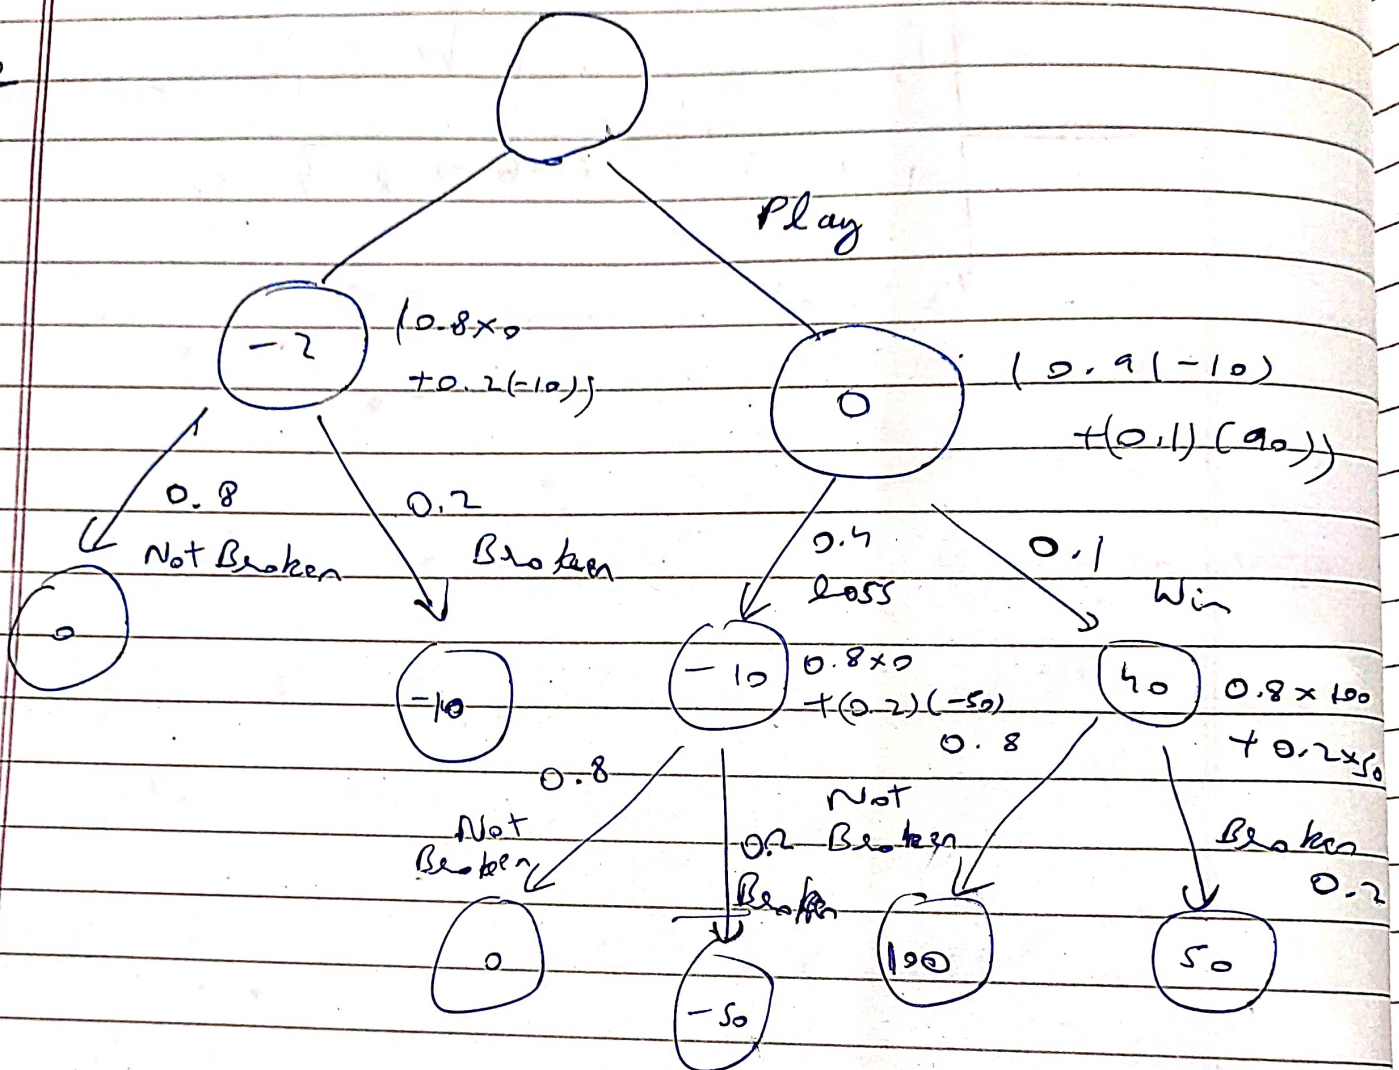
0.10128571428571428



0.1 learning rate works best with 30 epochs, with hidden layers as [256, 128, 64, 32], batch size = $\text{len}(\text{X_train})//20$, activation = ReLU. This worked best, as it is the optimum step the model takes after every step. So, it is able to reach the global minima faster and efficiently. It does not get stuck at the local minima like 0.001 and does not pass fast like 1.

[]:

Q.3

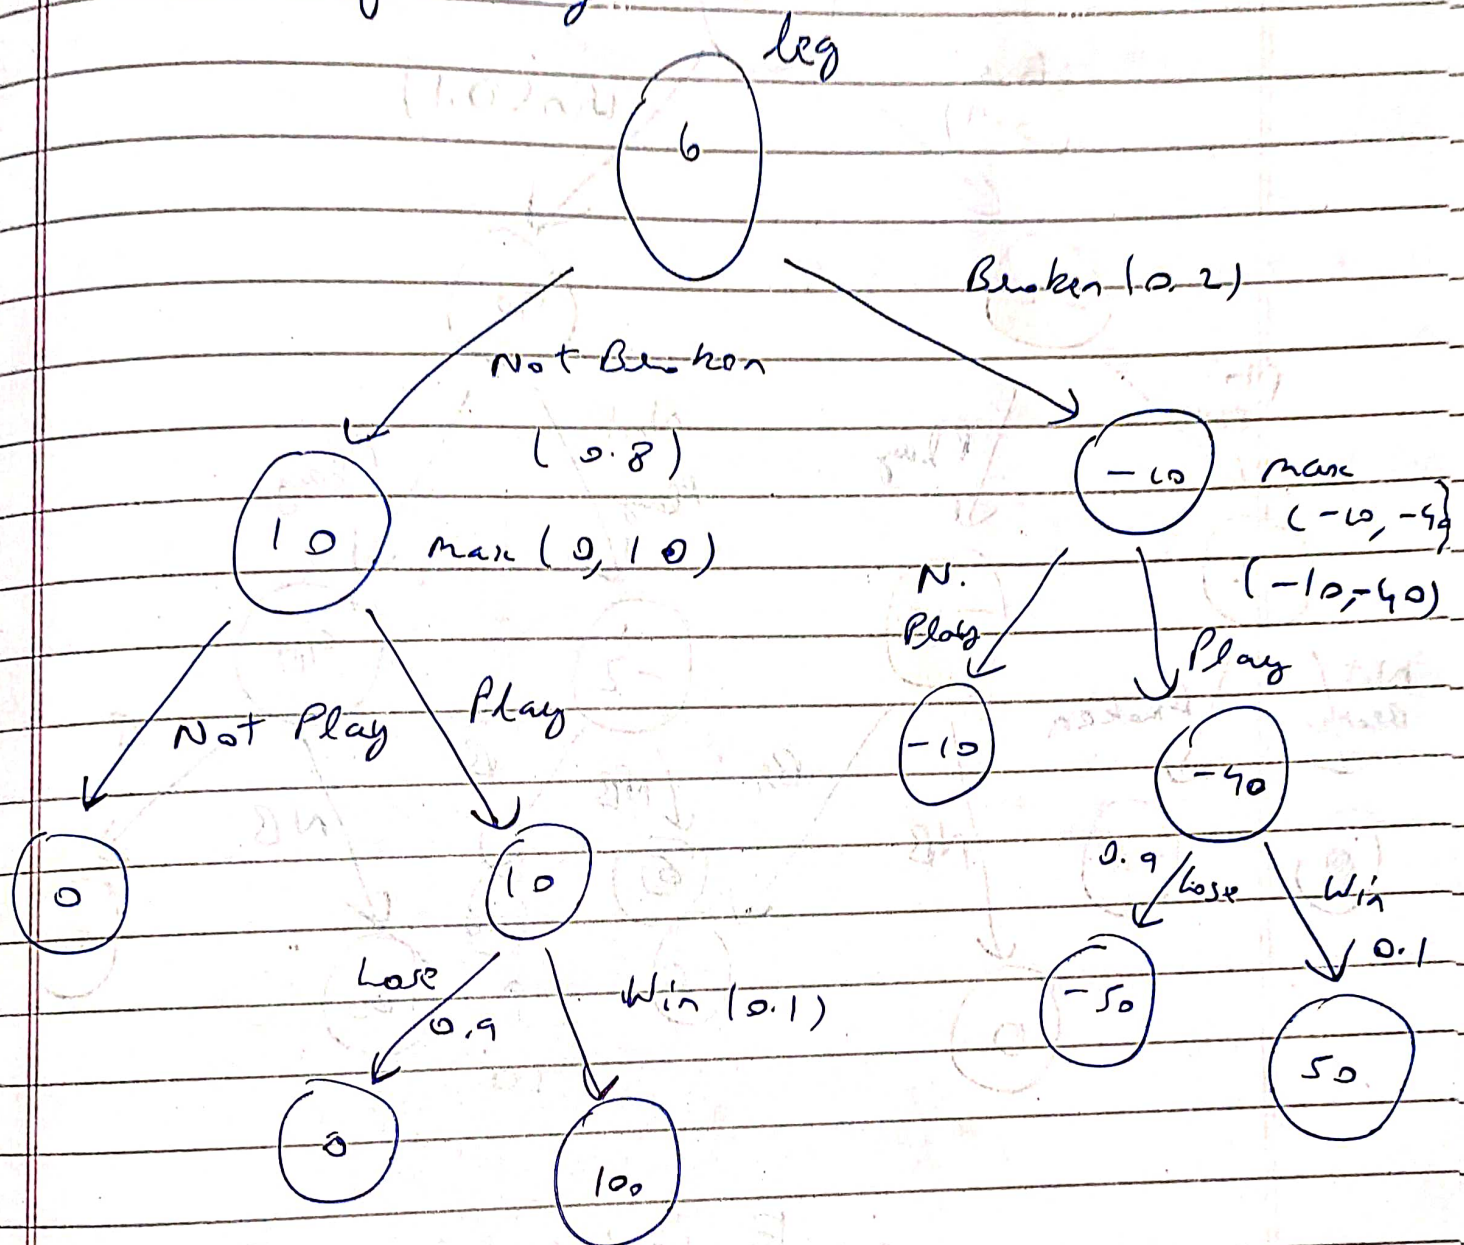
1/2

$$E(\text{play}) = 0$$

$$E(\text{not play}) = -2$$

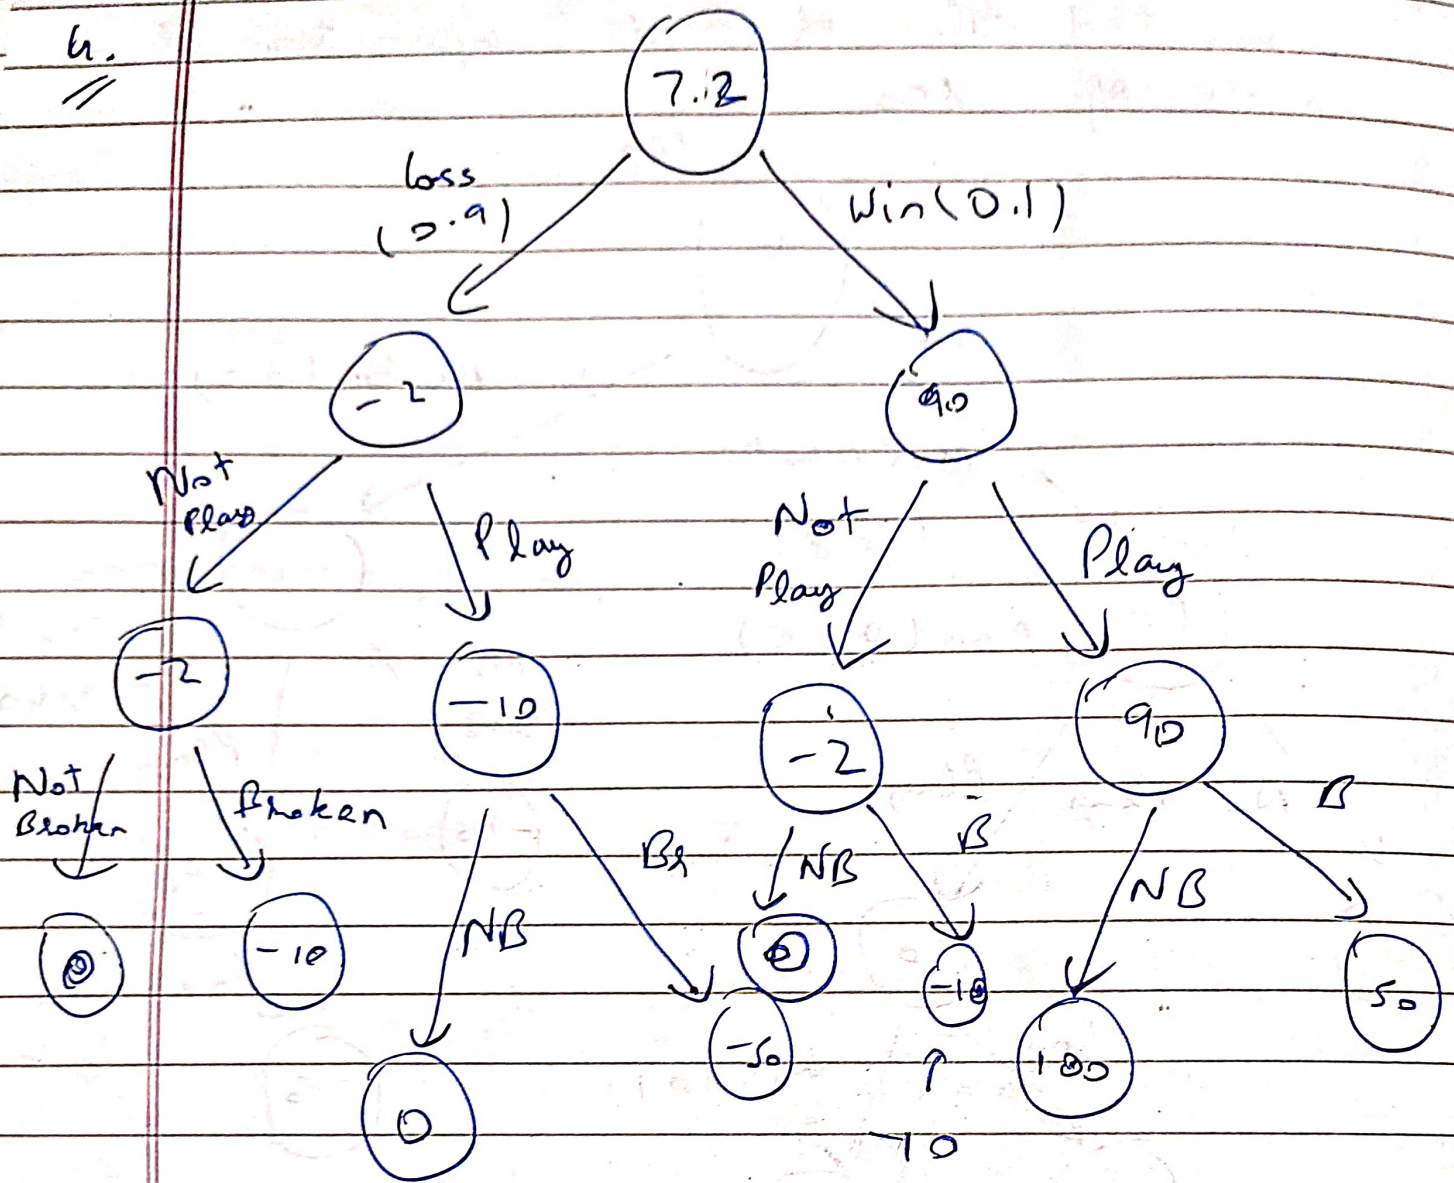
\therefore We should play

3. Expected value of perfect information about state of leg



\therefore Expected value of perfect information about state of leg
= 6

h.



\therefore Expected value = 72

Q5

Yes, on putting the win branch after the Broken branch.

Q2

Given $x \in \{0, 1\}^d$

We have $x = [x_1, x_2, \dots, x_d]$
where $x_i \in \{0, 1\}$

Now, we know that the conjunction operator can be modelled by a single neuron.

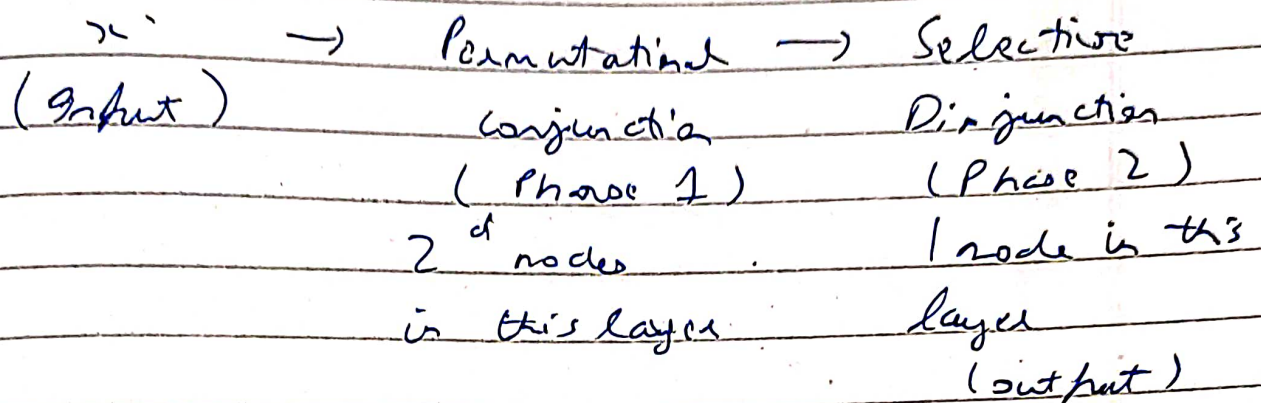
\Rightarrow We can model all possible 2^d conjunctions as "possible" inputs thus making size of hidden layer as 2^d .

Now, since we have all possible conjunctions for a certain d -dimensional input we can simply take the disjunction of all the codes we want to include in the final function. $\{ \}$ Reverse K-Map 3

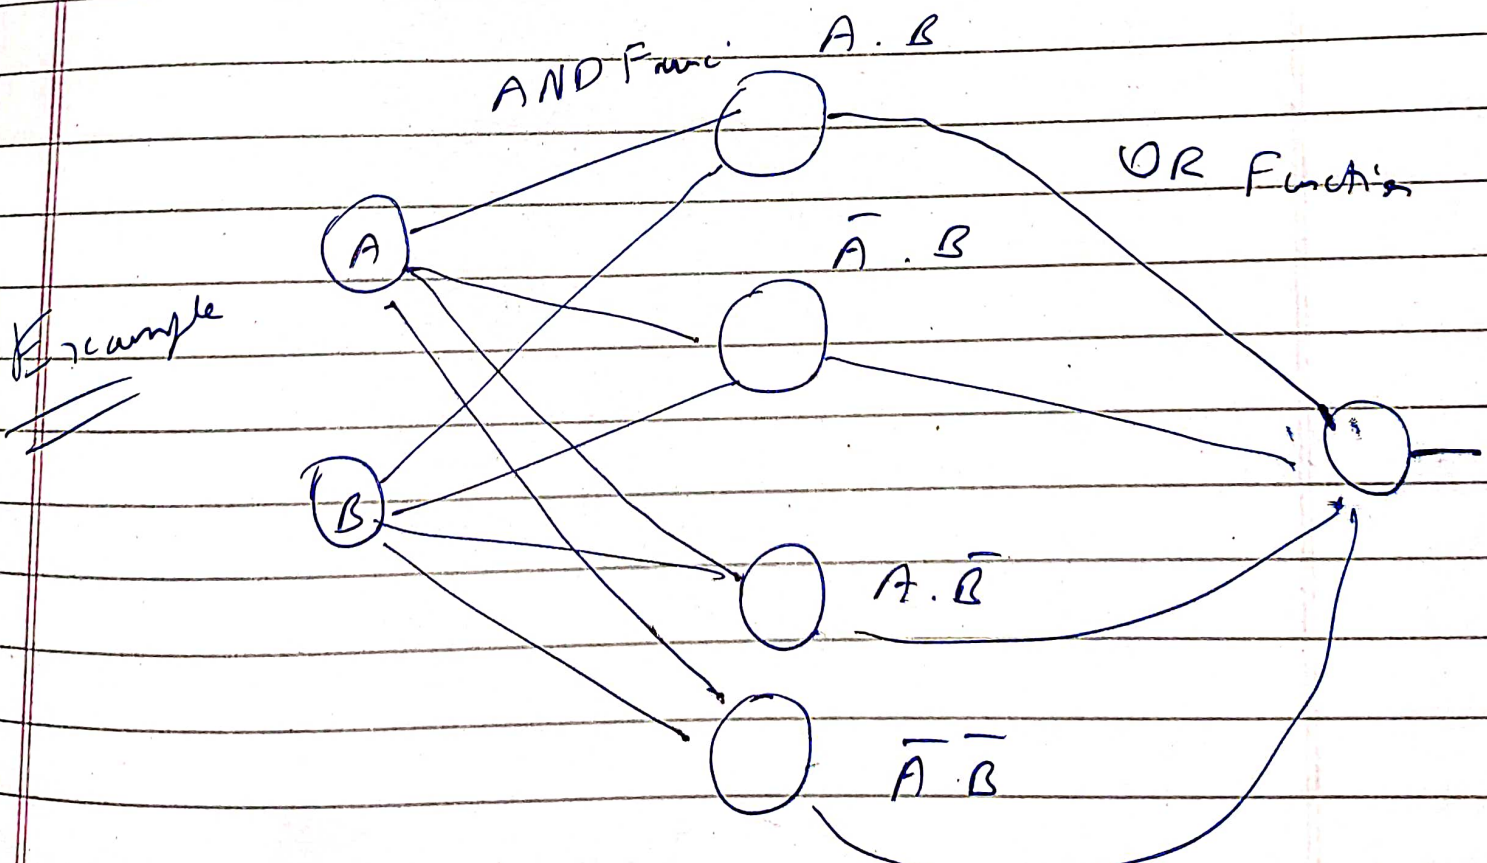
Since, a disjunction can be modelled by a single neuron,

\therefore The second layer can act as a selective disjunction

Hence, we split the listing of all possible functions in two phases, each of which we a single layer



Hence, a 2-layered NN can approximately model any function. When dealing with ~~to~~ boolean values, a hard threshold threshold activation can be used.



Q3

$$P(Y=1 | X) = \frac{e^{\beta_1 x_1 + \beta_2 x_2}}{1 + e^{\beta_1 x_1 + \beta_2 x_2}}$$

$$= \frac{1}{1 + \exp(-\beta_1 x_1 - \beta_2 x_2)}$$

$$\downarrow$$

$$\sigma(\beta_1 x_1 + \beta_2 x_2)$$

$$X = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

$$W = [\beta_1, \beta_2]$$

$$Z = WX = [\beta_1 x_1 + \beta_2 x_2]$$

